

harikrishnadevhw1-pytorch

August 28, 2023

1 HW1 (15 Points):

Required Submissions:

You have to submit two files for this part of the HW 1. FirstNameLastName_Hw1a.ipynb (colab notebook) 2. FirstNameLastName_Hw1a.pdf pdf file.

The pdf file is just the pdf version of colab notebook.

```
[11]: import torch
import time
import numpy as np
import matplotlib.pyplot as plt
```

```
[12]: torch.set_printoptions(precision=4, sci_mode=False)
```

Q1 : Create Tensor (1 Point) Create a torch Tensor of shape (5, 3) which is filled with zeros. Modify the tensor to set element (0, 2) to 10 and element (2, 0) to 100.

```
[13]: tensor = torch.zeros((5, 3))

tensor[0, 2] = 10
tensor[2, 0] = 100

print(tensor)
```

```
tensor([[ 0.,  0., 10.],
        [ 0.,  0.,  0.],
        [100.,  0.,  0.],
        [ 0.,  0.,  0.],
        [ 0.,  0.,  0.]])
```

2 Q2: Reshape tensor (1 Point)

You have following tensor as input:

```
x=torch.tensor([0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18,
19, 20, 21, 22, 23])
```

Using only reshaping functions (like view, reshape, transpose, permute), you need to get at the following tensor as output:

```
tensor([[ 0,  4,  8, 12, 16, 20],
        [ 1,  5,  9, 13, 17, 21],
        [ 2,  6, 10, 14, 18, 22],
        [ 3,  7, 11, 15, 19, 23]])
```

```
[14]: x=torch.tensor([0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23])
```

```
[15]: reshaped = x.view(6, 4)
      output = reshaped.transpose(0, 1)
      print(output)
```

```
tensor([[ 0,  4,  8, 12, 16, 20],
        [ 1,  5,  9, 13, 17, 21],
        [ 2,  6, 10, 14, 18, 22],
        [ 3,  7, 11, 15, 19, 23]])
```

3 Question 3. Speedtest for vectorization - 2 Points

Your goal is to measure the speed of linear algebra operations for different levels of vectorization.

1. Construct two matrices A and B with Gaussian random entries of size 1024×1024 .
2. Compute $C = AB$ using matrix-matrix operations and report the time. (Hint: Use torch.mm)
3. Compute $C = AB$, treating A as a matrix but computing the result for each column of B one at a time. Report the time. (hint use torch.mv inside a for loop)
4. Compute $C = AB$, treating A and B as collections of vectors. Report the time. (Hint: use torch.dot inside nested for loop)

```
[16]: ## Solution 1
      torch.manual_seed(42) # do not change this
      A = torch.randn(1024, 1024)
      B = torch.randn(1024, 1024)
```

```
[17]: ## Solution 2
      start=time.time()

      C = torch.mm(A, B)

      print("Matrix by matrix: " + str(time.time()-start) + " seconds")
```

Matrix by matrix: 0.05514264106750488 seconds

```
[18]: ## Solution 3
      C= torch.empty(1024,1024)
      start = time.time()
```

```

for i in range(1024):
    C[:, i] = torch.mv(A, B[:, i])

print("Matrix by vector: " + str(time.time()-start) + " seconds")

```

Matrix by vector: 0.29326939582824707 seconds

```

[19]: ## Solution 4
C= torch.empty(1024,1024)
start = time.time()

for i in range(1024):
    for j in range(1024):
        C[i, j] = torch.dot(A[i], B[j])

print("vector by vector: " + str(time.time()-start) + " seconds")

```

vector by vector: 15.7226083278656 seconds

4 Question 4 : Redo Question 3 by using GPU - 2 Point

Using GPUs

How to use GPUs in Google Colab In Google Colab – Go to Runtime Tab at top – select change runtime type – for hardware accelartor choose GPU

```

[20]: # Check if GPU is available
device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')
print(device)

```

cuda:0

```

[21]: ## Solution 1
torch.manual_seed(42)
A= torch.randn((1024, 1024),device=device)
B= torch.randn((1024, 1024),device=device)

```

```

[22]: ## Solution 2
start=time.time()

C = torch.mm(A, B)

print("Matrix by matrix: " + str(time.time()-start) + " seconds")

```

Matrix by matrix: 2.789036273956299 seconds

```

[23]: ## Solution 3
C= torch.empty(1024,1024, device = device)

```

```

start = time.time()

for i in range(1024):
    C[:, i] = torch.mv(A, B[:, i])

print("Matrix by vector: " + str(time.time()-start) + " seconds")

```

Matrix by vector: 0.05862569808959961 seconds

```

[24]: ## Solution 4
C= torch.empty(1024,1024, device = device)
start = time.time()

for i in range(1024):
    for j in range(1024):
        C[i, j] = torch.dot(A[i], B[j])

print("vector by vector: " + str(time.time()-start) + " seconds")

```

vector by vector: 36.108577728271484 seconds

5 Question 5. Memory efficient computation - 2 Points

We want to compute $C \leftarrow A \cdot B + C$, where A, B and C are all matrices. Implement this in the most memory efficient manner. Pay attention to the following two things:

1. Do not allocate new memory for the new value of C .
2. Do not allocate new memory for intermediate results if possible. Hint: If you implement this correctly the memory location of C given by `id(C)` will be same in both the cells below.

```

[25]: A= torch.randn((1000, 1000),device=device)
B= torch.randn((1000,1000),device=device)
C= torch.randn((1000, 1000),device=device)
print(id(C))

```

135463344210064

```

[26]: C.addmm_(A, B)

print(id(C))

```

135463344210064

#Question 6. Broadcast Operations - 2 Points

In order to perform polynomial fitting we want to compute a design matrix A with

$$A_{ij} = x_i^j$$

Our goal is to implement this **without a single for loop** entirely using vectorization and broadcast. Here $1 \leq j \leq 3$ and $x = \{1, 2, 3, 4, 5\}$. Implement code that generates following A matrix

$$\begin{bmatrix} 1 & 1 & 1 \\ 2 & 4 & 8 \\ 3 & 9 & 27 \\ 4 & 16 & 64 \\ 5 & 25 & 125 \end{bmatrix}$$

```
[27]: x = torch.tensor([1, 2, 3, 4, 5], dtype=torch.float32)
powers = torch.arange(1, 4, dtype=torch.float32)
x_powers = x.view(-1,1) ** powers
print(x_powers)
```

```
tensor([[ 1.,  1.,  1.],
        [ 2.,  4.,  8.],
        [ 3.,  9., 27.],
        [ 4., 16., 64.],
        [ 5., 25., 125.]])
```

6 Q7 Image Classification using - Pixel Similarity - 5 Points

6.1 Import libraries

```
[28]: # import libraries
import torchvision
import torchvision.transforms as transforms
from pathlib import Path
import cv2 as cv
from google.colab.patches import cv2_imshow # for image display
import pandas as pd
```

6.2 Mount Google Drive

```
[29]: # mount google drive
from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

6.3 Path to Download Data

Make sure you change the Path to where you want to save data. In the code below - **data/datasets** is the folder name in my google drive. You can change this to appropriate folder for your drive. For example you may want to save data to BUAN6341/HW1/Data. In this case the below code should be modified to :

```
data_folder = Path('/content/drive/MyDrive/BUAN6341/HW2/Data')
```

```
[30]: data_folder = Path('/content/drive/MyDrive/Colab Notebooks/BUAN 6382/Data')
```

6.4 Download MNIST training Data

The MNIST is a dataset of handwritten digits, available from this page <http://yann.lecun.com/exdb/mnist/>. It has a training set of 60,000 examples, and a test set of 10,000 examples. In this task we will use images for digits 3 and 7.

```
[31]: trainset = torchvision.datasets.MNIST(root = data_folder,
                                           train = True,
                                           download = True,
                                           )
```

6.5 Subset of images

trainset.data has images and trainset.targets has the labels. Below we will create a Boolean mask for digits 3 and 7. We will use this mask to extract subset of images.

```
[32]: # get the mask
idx3 = trainset.targets==3
idx7 = trainset.targets==7
```

```
[33]: idx3 # This take the value of True wherever the label was 3 and False otherwise
```

```
[33]: tensor([False, False, False, ..., False, False, False])
```

```
[34]: threes = trainset.data[idx3] # use Boolean mask to extract images for digit 3
      ↪ three
sevens = trainset.data[idx7] # use Boolean mask to extract images for digit 7
      ↪ seven
```

```
[35]: threes.shape, sevens.shape
```

```
[35]: (torch.Size([6131, 28, 28]), torch.Size([6265, 28, 28]))
```

```
[36]: type(threes[0])
```

```
[36]: torch.Tensor
```

```
[37]: # get a sample image
sample_img_3 = threes[0]
```

```
[38]: sample_img_3.shape
```

```
[38]: torch.Size([28, 28])
```

6.5.1 Task1

Reshape the image so that it has following shape: 28, 28, 1. Then convert the tensor to numpy array.

```
[39]: sample_img_3_numpy = sample_img_3.unsqueeze(2).numpy()
```

```
[40]: # let us look at the image
# the above steps were needed because the function cv2_imshow needs a three-
    ↪dimensional numpy array of the
# shape H x W x C (height, width, number of channels). The number of channels-
    ↪for black and white images is 1.
cv2_imshow(sample_img_3_numpy)
```



6.5.2 Task2

Repeat the above steps to show the first image for digit 7. - Get the first image. - Reshape the image to 28, 28, 1 - Convert the tensor to numpy array

```
[41]: sample_img_7_numpy = sevens[0].unsqueeze(2).numpy()
```

```
[42]: cv2_imshow(sample_img_7_numpy)
```



6.6 Visualizing images as pixel values

- In a computer, images are stored in the form of matrices. The numbers in the matrix are called “pixel values.” - These pixel values show how bright each pixel is. - 0 means “black,” and 255 means “white.” - The matrix of numbers is called the channel, and there is only one channel in a grayscale image. - For color images, we have three channels (RGB - Red, Green Blue). That is we need three matrices to store color images.

```
[43]: df = pd.DataFrame(sample_img_3[5:24, 7:22].numpy())
df.style.set_properties(**{'font-size': '8pt'})
```

```
[43]: <pandas.io.formats.style.Styler at 0x7b32c9e63730>
```

- In the matrix below we have flipped color just for illustration (We are using 0 for white instead of black).

```
[44]: df.style.set_properties(**{'font-size': '8pt'}).background_gradient('Greys')
```

```
[44]: <pandas.io.formats.style.Styler at 0x7b32c9e63c70>
```

6.6.1 Task3

Convert threes and sevens generated earlier to float and and take the mean along dimension 0.

Hint(You can use `tensor.float()` and then take the mean along dimesnion 0.

Here we are generating an average image by taking the average of each pixel across images.

```
[45]: mean_threes = threes.mean(dim=0, dtype=torch.float32)
      mean_sevens = sevens.mean(dim=0, dtype=torch.float32)
```

```
[46]: mean_threes.shape
      mean_sevens.shape
```

```
[46]: torch.Size([28, 28])
```

```
[47]: cv2_imshow(mean_threes.unsqueeze(dim=2).numpy())
      # print()
      cv2_imshow(mean_sevens.unsqueeze(dim=2).numpy())
```



6.7 Prediction for validation images

We will calculate the distance of each validation image from the `mean_threes` and `mean_sevens`. If the validation image is closest to `mean_threes` then we will predict 1 else we will predict 0.

6.8 Get Valid Dataset

```
[48]: validset = torchvision.datasets.MNIST(root = data_folder,
      train = False,
      download = True,
      )
```



```
[49]: # get the mask
      idx3 = validset.targets==3
      idx7 = validset.targets==7
```

```
[50]: validset.data.shape
```

```
[50]: torch.Size([10000, 28, 28])
```

```
[51]: # get images and labels using the mask
      valid_data_3_7 = validset.data[idx3+idx7]
      valid_targets_3_7 = validset.targets[idx3 + idx7]
```

```
[52]: # check the shape of the inputs (images)
      # we have 2038 images with size 28 x 28.
      valid_data_3_7.shape
```

```
[52]: torch.Size([2038, 28, 28])
```

```
[53]: # labels for valid dataset
      valid_targets_3_7
```

```
[53]: tensor([7, 7, 3, ..., 3, 7, 3])
```

```
[54]: # change the label to 1 where the label was 3
      valid_targets_3_7[valid_targets_3_7==3] = 1
```

```
[55]: # change the label to 0 where the label was 7
      valid_targets_3_7[valid_targets_3_7==7] = 0
```

```
[56]: valid_targets_3_7.unique()
```

```
[56]: tensor([0, 1])
```

```
[57]: valid_targets_3_7
```

```
[57]: tensor([0, 0, 1, ..., 1, 0, 1])
```

```
[58]: valid_data_3_7.shape
```

```
[58]: torch.Size([2038, 28, 28])
```

6.9 Distance between images

6.9.1 Task4

Write a function to calculate the distance between two images a and b. - Calculate the difference between corresponding pixels of two images. - Calculate square of differences. - Take the

mean (take the mean across last two dimensions (-1, -2) of the square of the differences and then take the square root.

```
[59]: def dist_images(a, b):  
        diff = a.float() - b.float()  
        squared_diff = diff ** 2  
        distance = torch.sqrt(squared_diff.mean((-1, -2)))  
        return distance
```

```
[60]: # calculate the distance of the each image in validation set from mean_threes  
        # because of the broadcasting we are able to use the above function which was  
        ↪written to  
        # calculate the distnace between two images  
        valid_3_dist = dist_images(valid_data_3_7, mean_threes)
```

```
[61]: # calculate the distance of the each image in validation set from mean_sevens  
        valid_7_dist = dist_images(valid_data_3_7, mean_sevens)
```

```
[62]: # if you have done everything right, you will observe following shapes.  
        valid_3_dist.shape, valid_7_dist.shape
```

```
[62]: (torch.Size([2038]), torch.Size([2038]))
```

6.10 Prediction Using Distance

6.10.1 Task5

```
[63]: # step 1 : generate the tensor filled with zeros that has the same shape as  
        ↪valid_target_3_7  
        prediction = torch.zeros_like(valid_targets_3_7)
```

```
[64]: # step 2: update the tensor prediction. Update the value to 1, when  
        ↪valid_3_dist < valid_7_dist.  
        # Hint (Generate a Boolean mask using the condition valid_3_dist < valid_7_dist,  
        ↪and use the mask to update value to 1 )  
  
        mask = valid_3_dist < valid_7_dist  
        prediction[mask] = 1
```

```
[65]: prediction
```

```
[65]: tensor([0, 0, 1, ..., 1, 0, 1])
```

6.11 Accuracy

```
[66]: correct = prediction == valid_targets_3_7
```

```
[67]: correct
```

```
[67]: tensor([True, True, True, ..., True, True, True])
```

```
[68]: 100 * correct.float().mean().item()
```

```
[68]: 96.6143250465393
```